

Gino Angelici

714-421-0820 | gino.angelici@gmail.com | [linkedin.com/in/gino-angelici](https://www.linkedin.com/in/gino-angelici) | github.com/ancientabacus | ginoangelici.com

EDUCATION

University of California, San Diego

Bachelor of Science in Data Science

Sep 2021 – Jun 2025

San Diego, CA

TECHNICAL SKILLS

Languages: Python, SQL (PostgreSQL, BigQuery, Snowflake), R, JavaScript/TypeScript, HTML/CSS, Linux

ML & AI: PyTorch, TensorFlow, Scikit-learn, XGBoost, Hugging Face Transformers, LangChain/LangGraph, pandas, NumPy, Statsmodels, SciPy, PySpark, Seaborn, matplotlib, Plotly, D3.js

Tools & Platforms: AWS (Lambda, EC2), Apache Spark, Airflow, Databricks, Snowflake, Docker, FastAPI, Git/GitHub, Weights & Biases, Tableau, Power BI, MongoDB

EXPERIENCE

Exponential Intelligence

Sep 2025 – Present

Machine Learning Intern

Irvine, CA

- Built an end-to-end LLM pipeline with prompt engineering and fine-tuned classification models to generate AI digital twins simulating survey responses across political subgroups; engineered iterative re-prompting strategies to correct model failure modes including social desirability bias and extreme response tendency.
- Developed a statistical validation framework using Jensen-Shannon Divergence and chi-square tests of homogeneity to measure distributional alignment between LLM-generated outputs and real survey respondents; identified domain-specific failure patterns across candidate evaluation and social attitude prediction tasks.

Halicioğlu Data Science Institute – Autonomous Vehicle Development

Sep 2024 – Aug 2025

ML Research Assistant

San Diego, CA

- Designed and deployed CNN-based perception and control models for autonomous vehicle navigation on low-power embedded hardware; integrated stereo and RGB-D cameras for real-time 3D scene understanding; leveraged UCSD's Supercomputer Center to train and fine-tune models at scale across 3 prototype configurations.
- Benchmarked custom CNN models against LLMs and LVMs; built MLOps workflows using Weights & Biases for experiment tracking and regression testing; led team to 1st place at Purdue University's national autonomous vehicle competition.

Scripps Institution of Oceanography

Sep 2022 – Sep 2024

Data Research Assistant

San Diego, CA

- Built Python ETL pipelines to ingest, clean, and analyze 70+ years of climate records across 60 coastal airports, processing millions of weather observations; applied sliding-window PCA, polynomial regression, and ridge regression to identify statistically significant correlations; contributed findings to a peer-reviewed publication update.

Deloitte

Mar 2023 – Aug 2023

Data Science Intern

San Diego, CA

- Trained and evaluated XGBoost and regularized regression models on 50,000+ HHS records; conducted feature importance analysis to surface behavioral and demographic predictors of substance abuse risk; communicated findings to senior leadership via interactive Seaborn and Plotly dashboards.

PROJECTS

Pokemon RAG AI Agent – LLM Knowledge Retrieval System

Mar 2026 – Present

- Built a RAG-powered conversational AI agent over a structured Pokemon knowledge base; implemented end-to-end document ingestion, text chunking, and vector embedding pipelines using LangChain and ChromaDB; designed semantic retrieval with dense vector search and prompt-engineered LLM responses for multi-turn question answering.

Amazon Purchase Dataset – Large-Scale ML Pipeline

Oct 2025 – Dec 2025

- Engineered end-to-end ML pipelines using AWS Lambda and Apache Spark to process and cleanse 16M+ rows of purchase data; implemented distributed PCA and feature engineering within BigQuery for scalable downstream model evaluation.

Reddit NLP Pipeline – Deep Learning Text Generation

Nov 2025 – Jan 2026

- Trained and evaluated a character-level deep learning model in TensorFlow on unstructured corpora from Reddit's top 200 communities; designed a benchmarking framework comparing architectural configurations against perplexity metrics; authored paper on domain-specific language pattern capture.