

Gino Angelici

gino.angelici@gmail.com | 714-421-0820 | linkedin.com/in/gino-angelici | ginoangelici.com

PROFESSIONAL SUMMARY

Data Scientist with a strong foundation in machine learning, LLM systems, and large-scale data pipelines, built through research and internship experience in tech-adjacent environments. Skilled in end-to-end model development from feature engineering and training to validation and deployment, with demonstrated ability to ship production-grade ML workflows and communicate findings to senior stakeholders.

EDUCATION

University of California, San Diego
Bachelor of Science in Data Science

San Diego, CA
Jun 2025

EXPERIENCE

Machine Learning Intern

Sep 2025 -Present

Exponential Intelligence

Irvine, CA

- Built an end-to-end LLM pipeline with prompt engineering and fine-tuned classification models to generate AI digital twins simulating survey responses across political subgroups; engineered iterative re-prompting strategies to correct failure modes including social desirability bias and extreme response tendency.
- Developed a statistical validation framework using Jensen-Shannon Divergence and chi-square tests of homogeneity to measure distributional alignment between LLM outputs and real survey respondents; identified domain-specific failure patterns across candidate evaluation and social attitude prediction tasks.

ML Research Assistant

Sep 2024 -Aug 2025

Halicioglu Data Science Institute, Autonomous Vehicle Development

San Diego, CA

- Designed and deployed CNN-based perception and control models for autonomous vehicle navigation on low-power embedded hardware; integrated stereo and RGB-D cameras for real-time 3D scene understanding; leveraged UCSD's Supercomputer Center to train and fine-tune models across 3 prototype configurations.
- Benchmarked CNN models against LLMs and LVMs; built MLOps workflows using Weights & Biases for experiment tracking and regression testing; led team to 1st place at Purdue University's national autonomous vehicle competition.

Data Research Assistant

Sep 2022 -Sep 2024

Scripps Institution of Oceanography

San Diego, CA

- Built Python ETL pipelines to ingest, clean, and analyze 70+ years of climate records across 60 airports; applied PCA, polynomial and ridge regression to identify statistically significant correlations; contributed findings to a peer-reviewed publication update.

Data Science Intern

Mar 2023 -Aug 2023

Deloitte

San Diego, CA

- Trained and evaluated XGBoost and regularized regression models on 50,000+ HHS records; conducted feature importance analysis to surface behavioral and demographic predictors of substance abuse risk; presented findings to senior leadership via Plotly dashboards.

PROJECTS

Pokemon RAG AI Agent

Mar 2026 -Present

LLM Knowledge Retrieval System

- Built a RAG-powered conversational AI agent over a structured knowledge base; implemented end-to-end document ingestion, text chunking, and vector embedding pipelines using LangChain and ChromaDB; designed semantic retrieval with dense vector search and prompt-engineered LLM responses for multi-turn QA.

Amazon Purchase Dataset

Oct 2025 -Dec 2025

Large-Scale ML Pipeline

- Engineered end-to-end ML pipelines using AWS Lambda and Apache Spark to process and cleanse 16M+ rows of purchase data; implemented distributed PCA and feature engineering within BigQuery for scalable downstream model evaluation.

Reddit NLP Pipeline

Nov 2025 -Jan 2026

Deep Learning Text Generation

- Trained and evaluated a character-level deep learning model in TensorFlow on Reddit's top 200 communities; designed a benchmarking framework comparing architectural configurations against perplexity metrics; authored paper on language pattern capture.

SKILLS

Languages: Python, SQL (PostgreSQL, BigQuery, Snowflake), R, JavaScript/TypeScript, HTML/CSS, Linux

ML & AI: PyTorch, TensorFlow, Scikit-learn, XGBoost, Hugging Face Transformers, LangChain/LangGraph, pandas, NumPy, PySpark, Statsmodels, SciPy, Seaborn, matplotlib, Plotly, D3.js

Tools & Platforms: AWS (Lambda, EC2), Apache Spark, Docker, Databricks, Airflow, Snowflake, Weights & Biases, FastAPI, Git/GitHub, Tableau, Power BI, MongoDB